

Policy AI Lab — What We Mean by "Learning"

Written follow-up to Dr. Muath's question · Meeting of 8 July 2026

Scope note: This note refers specifically to the Policy AI Lab solution and should not be generalised to other solutions.

1. Summary

Thank you for raising the question in the session, and for setting out the framing so clearly. To confirm it in writing: for this solution we are not training or fine-tuning a large language model (LLM). The core engine is a pre-trained, off-the-shelf (open-source) model that we use without modifying its weights. What we refer to as "learning" is, consistent with the framing you outlined, a combination of in-context (few-shot) prompting, retrieval over the uploaded document corpus, and human-in-the-loop refinement of the flows and prompts. Quality over time is maintained through an evaluation-based approach to model selection.

2. How the Policy AI Lab agents work

2.1 Few-shot / in-context prompting

In line with the few-shot approach you described, the agents are guided with curated examples rather than trained. In practice this means we supply reference examples at the relevant steps so the agents follow the intended pattern, and we build up context across steps where needed. The historical BCG policy documents (with version control) that you kindly offered to share are exactly the input we need: we will extract the process and the pattern — not the published numbers — and turn them into the agents' few-shot examples. This is where your reference material will most directly improve output quality.

2.2 Retrieval over uploaded documents

When we say the system "learns from your documents", we mean it retrieves from the uploaded policy corpus at inference time and uses that content to ground its output. The model reads the relevant material on demand; its weights are not updated. As more high-quality, correctly-governed documents are made available through the secure sandbox, the grounding improves — this is retrieval, not training.

2.3 Human-in-the-loop refinement

Improvement comes from people iterating on the agent flow and the agent prompts. When an output needs correction, the fix is made by adjusting the flow, the agent flow, or the specific agent prompt — not by the model changing itself. Where useful, a new agent prompt can be refined further through systematic prompt optimisation; we treat that as an option rather than routine operation.

3. Quality assurance and model selection

For measuring and sustaining quality, our intention is to align with the evaluation-based approach your team already applies — building evaluation pipelines that compare outputs against a baseline and score them (including LLM-as-judge scoring anchored on human-validated evaluation sets), so that the strongest available model can be selected objectively over time. Framing it precisely: improvements in capability come from selecting and adopting stronger models as the open-source ecosystem advances — not from us training a model. Working this way is a deliberate design choice for this solution: it keeps the system flexible, avoids locking into any single model, and lets us adopt stronger models as they emerge.

4. Consistent message for stakeholders

For communication with stakeholders, we suggest a single consistent message:

"Policy AI Lab improves through better few-shot examples and prompts, through retrieval over governed documents, and through evaluation-based selection of the strongest available model. The core LLM is a pre-trained, off-the-shelf engine that we do not train."